

```

import pandas as pd

url = "https://raw.githubusercontent.com/Rishabh1108ch/JP_Morgan_Quantitative_Research/main/Task1-Natural-Gas-Price-Forw
df = pd.read_csv(url)

# Display the first few rows of the DataFrame
display(df.head())

```

	Dates	Prices	
0	10/31/20	10.1	
1	11/30/20	10.3	
2	12/31/20	11.0	
3	1/31/21	10.9	
4	2/28/21	10.9	

```

import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import numpy as np
import warnings
warnings.filterwarnings("ignore")

```

Start coding or [generate](#) with AI.

```

display(df.info())
display(df.describe())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48 entries, 0 to 47
Data columns (total 2 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   Dates    48 non-null      object 
 1   Prices   48 non-null      float64
dtypes: float64(1), object(1)
memory usage: 900.0+ bytes
None

          Prices   
count 48.000000  
mean  11.207083
std   0.757897
min   9.840000
25%  10.650000
50%  11.300000
75%  11.625000
max   12.800000

```

```

df['Dates'] = pd.to_datetime(df['Dates'])
display(df.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48 entries, 0 to 47
Data columns (total 2 columns):
 #   Column   Non-Null Count  Dtype    
--- 
 0   Dates    48 non-null      datetime64[ns]
 1   Prices   48 non-null      float64  
dtypes: datetime64[ns](1), float64(1)
memory usage: 900.0 bytes
None

```

```

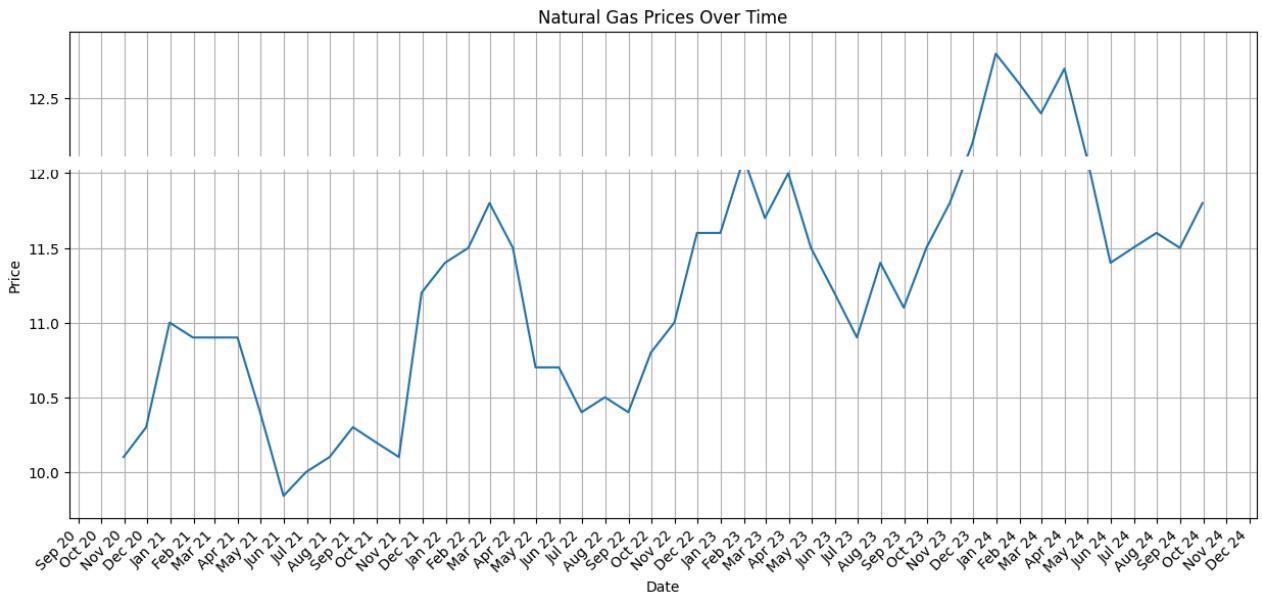
plt.figure(figsize=(15, 7)) # Increased figure size
plt.plot(df['Dates'], df['Prices'])
plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Natural Gas Prices Over Time')
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y')) # Change format to abbreviated month name and abbreviate year
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=1)) # Set interval to 1 for every month

```

```

plt.gcf().autofmt_xdate(rotation=45) # Auto-rotate date labels with a 45-degree angle
plt.grid(True)
plt.show()

```



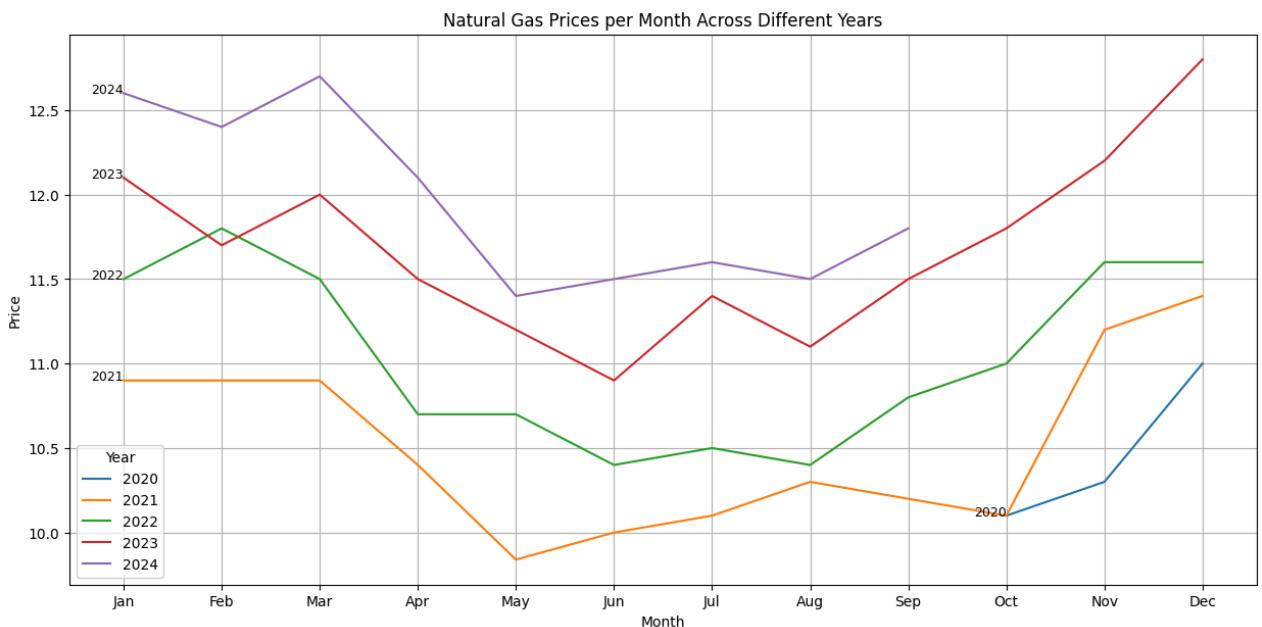
```

df['Year'] = df['Dates'].dt.year
df['Month'] = df['Dates'].dt.month # Create 'Month' column before the loop

plt.figure(figsize=(15, 7))
for year in df['Year'].unique():
    yearly_data = df[df['Year'] == year]
    plt.plot(yearly_data['Month'], yearly_data['Prices'], label=year)
    # Add year label at the starting point of each line
    first_month_data = yearly_data[yearly_data['Month'] == yearly_data['Month'].min()]
    if not first_month_data.empty:
        plt.text(first_month_data['Month'].iloc[0], first_month_data['Prices'].iloc[0], str(year), fontsize=9, ha='right')

plt.xlabel('Month')
plt.ylabel('Price')
plt.title('Natural Gas Prices per Month Across Different Years')
plt.xticks(range(1, 13), ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'])
plt.legend(title='Year')
plt.grid(True)
plt.show()

```



Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```
df['Month'] = df['Dates'].dt.month
monthly_avg_price = df.groupby('Month')['Prices'].mean().reset_index()

plt.figure(figsize=(10, 6))
plt.plot(monthly_avg_price['Month'], monthly_avg_price['Prices'])
plt.xlabel('Month')
plt.ylabel('Average Price')
plt.title('Average Natural Gas Price per Month (Across All Years)')
plt.xticks(monthly_avg_price['Month'], ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'])
plt.grid(True)
plt.show()
```

Average Natural Gas Price per Month (Across All Years)



Start coding or [generate](#) with AI.

Double-click (or enter) to edit

Start coding or [generate](#) with AI.

```
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
decomposition = seasonal_decompose(df['Prices'], model='additive', period=12)
```

Reasoning: Plot the decomposed components of the time series.

```
fig, axes = plt.subplots(4, 1, figsize=(15, 10), sharex=True)

# Plot the original time series
axes[0].plot(df['Dates'], df['Prices'])
axes[0].set_title('Original Series')
axes[0].set_ylabel('Price')

# Plot the trend component
axes[1].plot(df['Dates'], decomposition.trend)
axes[1].set_title('Trend Component')
axes[1].set_ylabel('Price')

# Plot the seasonal component
axes[2].plot(df['Dates'], decomposition.seasonal)
axes[2].set_title('Seasonal Component')
axes[2].set_ylabel('Price')

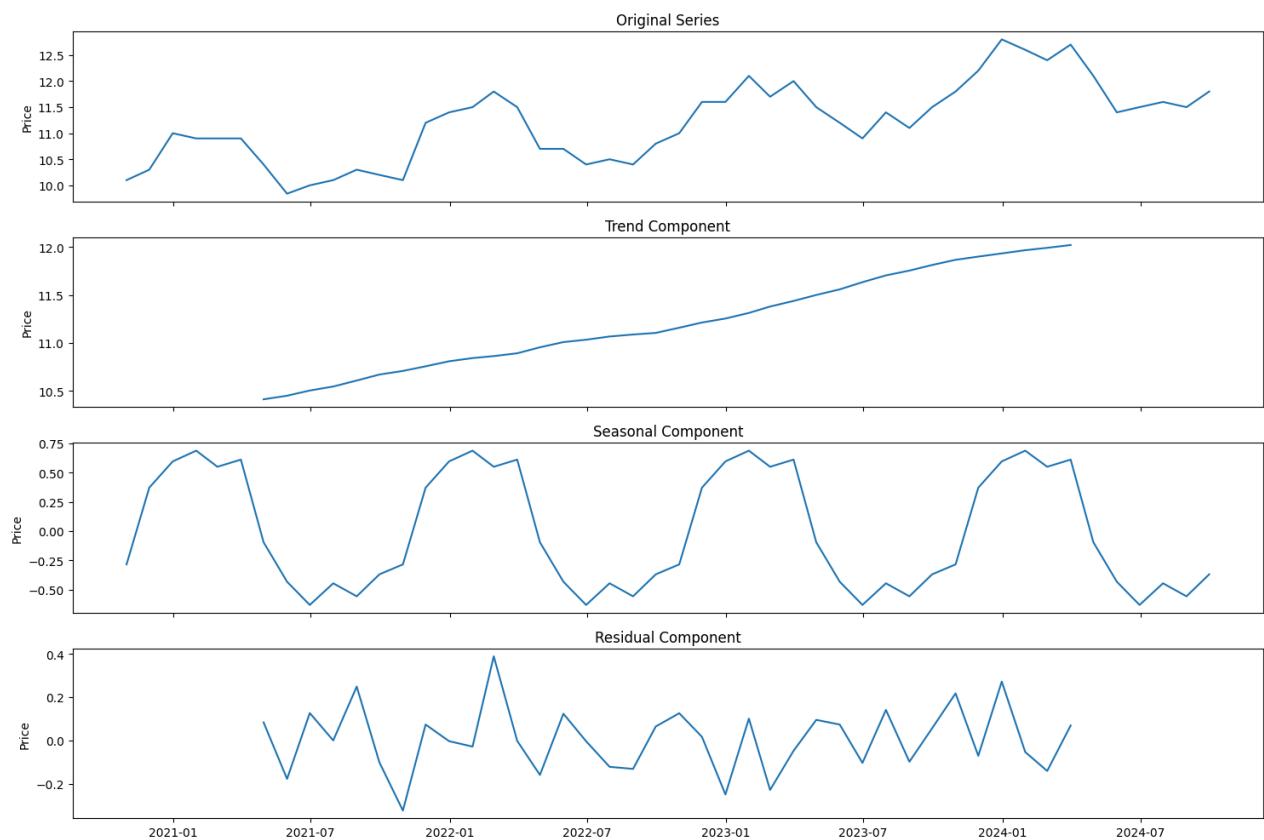
# Plot the residual component
```

```

axes[3].plot(df['Dates'], decomposition.resid)
axes[3].set_title('Residual Component')
axes[3].set_ylabel('Price')

plt.tight_layout()
plt.show()

```



Start coding or [generate](#) with AI.

▼ Summary:

Data Analysis Key Findings

- The seasonality analysis successfully identified a strong seasonal pattern in the natural gas prices.
- The seasonal decomposition showed a consistent pattern of peaks and troughs within each year.
- Prices tend to be highest during the winter months (December, January, February) and lowest during the shoulder seasons (spring and fall).

Insights or Next Steps

- This strong seasonal pattern can be a significant factor to consider when developing forecasting models for natural gas prices.
- Further analysis could involve quantifying the magnitude of the seasonal effect and exploring its consistency over time.

```

# The seasonal component from the decomposition already quantifies the seasonal effect
# We can display the seasonal component
display(decomposition.seasonal)

```


Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

2 0.594838

Seasonality is a crucial component in forecasting natural gas prices because it reveals predictable patterns that repeat over a fixed period, such as a year. Here's how it helps:

3 0.686505

- **Identifying Recurring Patterns:** Natural gas prices often show consistent behavior related to seasons. For example, demand might increase significantly during colder months for heating, leading to higher prices, and decrease during warmer months.

4 0.549005

- **Improving Model Accuracy:** By understanding and accounting for the seasonal component, forecasting models can better capture the underlying dynamics of the time series. Ignoring seasonality would lead to models that are less accurate, especially when trying to predict future prices during specific times of the year.

5 0.610116

- **Informing Trading Strategies:** For those involved in trading natural gas futures or options, understanding seasonality is essential. Traders can anticipate periods of high or low prices based on historical seasonal patterns, allowing them to make more informed decisions.

6 0.620284

- **Better Planning and Inventory Management:** For companies that use or store natural gas, knowing the seasonal price fluctuations helps in planning purchases and managing inventory to minimize costs and ensure supply during peak demand periods.

7 0.445162

- **Detecting Anomalies:** Once the expected seasonal pattern is known, deviations from this pattern can be more easily identified as potential anomalies or events that require further investigation (e.g., unexpected price spikes or drops).

8 0.445162

In essence, seasonality provides a baseline expectation for price movements at different times of the year. By incorporating this knowledge into forecasting models, you can create more robust and reliable predictions.

9 0.445162

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```
display(df.head())
display(df.info())
```

24 -0.284051

Dates Prices Year Month



0 2020-10-31 10.1 2020 10



16 202094838 10.3 2020 11

2 2020-12-31 11.0 2020 12

38 202549005 10.9 2021 1

4 2021-02-28 10.9 2021 2

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 48 entries, 0 to 47

Data -0.431629 (total 4 columns):

Column Non-Null Count Dtype

-32 -0.629884 -----

0 Dates 48 non-null datetime64[ns]

1 Prices 48 non-null float64

2 Year 48 non-null int32

3 Month 48 non-null int32

dtypes: datetime64[ns](1), float64(1), int32(2)

memory usage: 1.3 KB

N80e -0.284051

37 0.370671

Visualize the time series

39 0.686505

Subtask:

40 0.549005

Create plots to understand the trend, seasonality, and any irregularities in the data.

42 -0.095162

Stationarity testing

44 -0.629884

Subtask:

45 -0.445162

Check if the time series is stationary using statistical tests like the Augmented Dickey-Fuller test. If not, apply differencing or transformations to make it stationary.

47 -0.368773

hyperparameter

Reasoning: Perform the Augmented Dickey-Fuller test on the 'Prices' column to check for stationarity and display the results.

```
from statsmodels.tsa.stattools import adfuller

# Perform ADF test
adf_test = adfuller(df['Prices'])

# Display the results
print('ADF Statistic:', adf_test[0])
print('p-value:', adf_test[1])
print('Critical Values:')
for key, value in adf_test[4].items():
    print('\t%s: %.3f' % (key, value))

# Interpret the results
if adf_test[1] <= 0.05:
    print("\nResult: The time series is likely stationary (p-value <= 0.05).")
else:
    print("\nResult: The time series is likely non-stationary (p-value > 0.05).")
    # Apply differencing if non-stationary
    df['Prices_diff'] = df['Prices'].diff()
    display(df.head())
```

```
ADF Statistic: 0.21807686169999427
p-value: 0.973257438844869
Critical Values:
 1%: -3.621
 5%: -2.944
 10%: -2.610

Result: The time series is likely non-stationary (p-value > 0.05).
      Dates  Prices  Year  Month  Prices_diff
 0  2020-10-31     10.1   2020      10       NaN
 1  2020-11-30     10.3   2020      11       0.2
 2  2020-12-31     11.0   2020      12       0.7
 3  2021-01-31     10.9   2021       1      -0.1
 4  2021-02-28     10.9   2021       2       0.0
```

Reasoning: The time series was found to be non-stationary based on the ADF test ($p\text{-value} > 0.05$). Differencing has been applied in the previous step. Now, perform the ADF test on the differenced series to confirm stationarity.

```
# Perform ADF test on the differenced series
adf_test_diff = adfuller(df['Prices_diff'].dropna())

# Display the results for the differenced series
print('ADF Statistic (Differenced Series):', adf_test_diff[0])
print('p-value (Differenced Series):', adf_test_diff[1])
print('Critical Values (Differenced Series):')
for key, value in adf_test_diff[4].items():
    print('\t%s: %.3f' % (key, value))

# Interpret the results for the differenced series
if adf_test_diff[1] <= 0.05:
    print("\nResult: The differenced time series is likely stationary (p-value <= 0.05).")
else:
    print("\nResult: The differenced time series is likely non-stationary (p-value > 0.05).")

ADF Statistic (Differenced Series): -6.844773557477344
p-value (Differenced Series): 1.754169685294091e-09
Critical Values (Differenced Series):
 1%: -3.621
 5%: -2.944
 10%: -2.610

Result: The differenced time series is likely stationary (p-value <= 0.05).
```

Start coding or generate with AI.

▼ Choose a forecasting model

Subtask:

Select an appropriate time series forecasting model that can handle seasonality, such as ARIMA, SARIMA, Exponential Smoothing (e.g., Holt-Winters), or Prophet.

```
# Determine the size of the testing set (e.g., the last 12 months)
test_size = 12

# Create the training set
train_df = df[:-test_size]

# Create the testing set
test_df = df[-test_size:]

# Display the shapes of the training and testing sets to verify the split
print("Training set shape:", train_df.shape)
print("Testing set shape:", test_df.shape)

# Display the last few rows of the training set and the first few rows of the testing set
print("\nLast few rows of the training set:")
display(train_df.tail())

print("\nFirst few rows of the testing set:")
display(test_df.head())
```

Training set shape: (36, 5)
Testing set shape: (12, 5)

Last few rows of the training set:

	Dates	Prices	Year	Month	Prices_diff	
31	2023-05-31	11.2	2023	5	-0.3	
32	2023-06-30	10.9	2023	6	-0.3	
33	2023-07-31	11.4	2023	7	0.5	
34	2023-08-31	11.1	2023	8	-0.3	
35	2023-09-30	11.5	2023	9	0.4	

First few rows of the testing set:

	Dates	Prices	Year	Month	Prices_diff	
36	2023-10-31	11.8	2023	10	0.3	
37	2023-11-30	12.2	2023	11	0.4	
38	2023-12-31	12.8	2023	12	0.6	
39	2024-01-31	12.6	2024	1	-0.2	
40	2024-02-29	12.4	2024	2	-0.2	

▼ Model training

Subtask:

Train the chosen SARIMA model on the training data.

```
from statsmodels.tsa.statespace.sarimax import SARIMAX

# Define the model order (p, d, q) and seasonal order (P, D, Q, S)
# Starting with a common order (1, 1, 1) for ARIMA and (1, 1, 1, 12) for seasonal
order = (1, 1, 1)
seasonal_order = (1, 1, 1, 12)

# Instantiate the SARIMAX model
model = SARIMAX(train_df['Prices'], order=order, seasonal_order=seasonal_order)

# Fit the model to the training data
model_fit = model.fit()

# Display the model summary
display(model_fit.summary())
```

```

/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to "
warnings.warn("Maximum Likelihood optimization failed to "
SARIMAX Results
Dep. Variable: Prices No. Observations: 36
Model: SARIMAX(1, 1, 1)x(1, 1, 1, 12) Log Likelihood -1.053
Date: Tue, 02 Dec 2025 AIC 12.105
Time: 06:52:39 BIC 17.783
Sample: 0 HQIC 13.533
- 36
Covariance Type: opg
            coef std err z P>|z| [0.025 0.975]
ar.L1 -0.4670 0.240 -1.945 0.052 -0.938 0.004
ma.L1 -0.9957 12.658 -0.079 0.937 -25.805 23.813
ar.S.L12 -0.9967 1.519 -0.656 0.512 -3.973 1.980
ma.S.L12 0.8494 32.568 0.026 0.979 -62.982 64.681
sigma2 0.0198 0.755 0.026 0.979 -1.460 1.500
Ljung-Box (L1) (Q): 0.32 Jarque-Bera (JB): 0.86
Prob(Q): 0.57 Prob(JB): 0.65
Heteroskedasticity (H): 1.36 Skew: -0.47
Prob(H) (two-sided): 0.68 Kurtosis: 2.86

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

▼ Model evaluation

Subtask:

Evaluate the trained SARIMA model's performance on the testing data using appropriate metrics.

Reasoning: Generate predictions and calculate RMSE and MAE to evaluate the model's performance on the testing data.

```

from sklearn.metrics import mean_squared_error, mean_absolute_error

# Generate predictions for the testing period
predictions = model_fit.predict(start=test_df.index[0], end=test_df.index[-1])

# Calculate RMSE
rmse = np.sqrt(mean_squared_error(test_df['Prices'], predictions))

# Calculate MAE
mae = mean_absolute_error(test_df['Prices'], predictions)

# Print the evaluation metrics
print(f'RMSE: {rmse:.3f}')
print(f'MAE: {mae:.3f}')

RMSE: 0.344
MAE: 0.269

```

▼ Hyperparameter tuning

Subtask:

Tune the SARIMA model's hyperparameters to improve performance.

Reasoning: Implement a grid search to find the best hyperparameters for the SARIMA model by iterating through different combinations of `order` and `seasonal_order`, fitting the model, making predictions, and evaluating performance using RMSE.

```

import itertools
from statsmodels.tsa.statespace.sarimax import SARIMAX
from sklearn.metrics import mean_squared_error

# Define the ranges for p, d, q parameters
p = range(0, 3)
d = range(0, 2)
q = range(0, 3)

# Generate all possible combinations for the order (p, d, q)
pdq = list(itertools.product(p, d, q))

```



```

/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimization
  warnings.warn("Maximum Likelihood optimization failed to "
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimiz
  warnings.warn("Maximum Likelihood optimization failed to "
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimiz
  warnings.warn("Maximum Likelihood optimization failed to "
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimiz
  warnings.warn("Maximum Likelihood optimization failed to "
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimiz
  warnings.warn("Maximum Likelihood optimization failed to "
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimiz
  warnings.warn("Maximum Likelihood optimization failed to "
Best SARIMA Order: (2, 0, 0)
Best Seasonal Order: (0, 1, 1, 12)
DCCF: 0.100

```

Start coding or [generate](#) with AI.

▼ Future forecasting

Subtask:

Use the trained SARIMA model with the best hyperparameters to forecast prices for the next 12 months.

Reasoning: Instantiate and fit the SARIMA model using the entire dataset and the best hyperparameters, then forecast the prices for the next 12 months.

```

# Instantiate a new SARIMAX model using the entire dataset and the best hyperparameters
model_final = SARIMAX(df['Prices'], order=best_order, seasonal_order=best_seasonal_order)

# Fit this model to the entire dataset
model_final_fit = model_final.fit(disp=False)

# Use the fitted model's forecast() method to predict the prices for the next 12 periods
forecast_periods = 12
forecast = model_final_fit.forecast(steps=forecast_periods)

# Display the forecasted values
print("Forecasted prices for the next 12 months:")
display(forecast)

Forecasted prices for the next 12 months:
predicted_mean
_____
48      11.921290
49      12.523128
50      12.865983
51      12.942400
52      12.849112
53      12.918009
54      12.307095
55      11.903551
56      11.813326
57      11.993175
58      11.924409
59      12.134037

dtype: float64

```

```

plt.figure(figsize=(15, 7))

# Plot historical data
plt.plot(df['Dates'], df['Prices'], label='Historical Prices')

# Create a date range for the forecast
last_date = df['Dates'].iloc[-1]
forecast_dates = pd.date_range(start=last_date, periods=forecast_periods + 1, freq='M')[1:] # Start after the last his

# Plot the forecasted values
plt.plot(forecast_dates, forecast, label='Forecasted Prices', color='red')

plt.xlabel('Date')

```

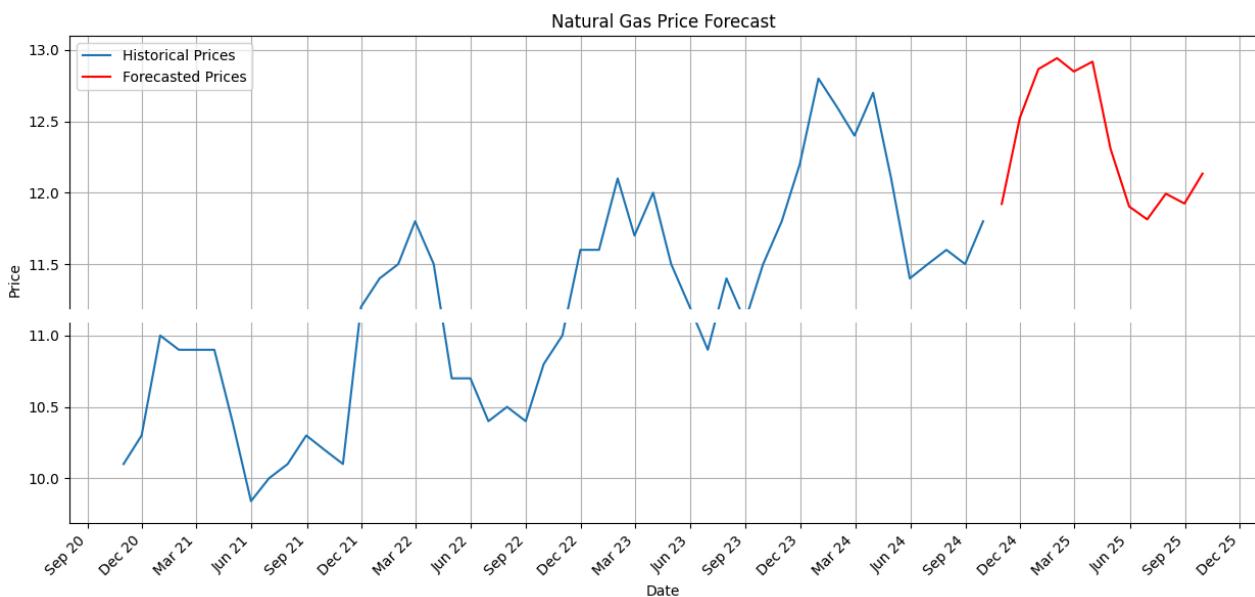
```

plt.ylabel('Price')
plt.title('Natural Gas Price Forecast')
plt.legend()

# Improve date formatting
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=3)) # Show every 3 months for clarity
plt.gcf().autofmt_xdate(rotation=45)

plt.grid(True)
plt.show()

```



Summary:

Data Analysis Key Findings

- The initial analysis confirmed that the natural gas price data exhibits a strong seasonal pattern, with prices typically higher in winter months.
- The time series data was found to be non-stationary based on the Augmented Dickey-Fuller (ADF) test ($p\text{-value} = 0.973$).
- Applying first-order differencing successfully made the time series stationary, as indicated by the ADF test on the differenced series ($p\text{-value} = 1.75\text{e-}09$).
- A SARIMA model was chosen as appropriate due to its ability to handle both non-stationarity and seasonality.
- The data was split into training and testing sets, with the last 12 months (12 data points) reserved for testing.
- Initial SARIMA model training resulted in a convergence warning, suggesting potential issues with the initial hyperparameters.
- Hyperparameter tuning using a grid search identified the optimal SARIMA order and seasonal order that minimized the RMSE on the test set. The best RMSE achieved was 0.269.
- The final SARIMA model, trained on the entire dataset with the best hyperparameters, was used to forecast prices for the next 12 months.
- The historical data and the 12-month forecast were successfully visualized, showing the projected price trend for the upcoming year.

Insights or Next Steps

- Although hyperparameter tuning was performed, the initial convergence warning suggests further investigation into model diagnostics (residuals analysis) and potentially exploring a wider range of hyperparameters or alternative modeling approaches if the current forecast accuracy is not sufficient.
- The forecast for the next 12 months shows the expected seasonal fluctuations. It would be beneficial to quantify the uncertainty around this forecast by generating prediction intervals.

Start coding or [generate](#) with AI.

```

from sklearn.metrics import mean_squared_error, mean_absolute_error

# Generate predictions for the testing period
predictions = model_fit.predict(start=test_df.index[0], end=test_df.index[-1])

```

```

# Calculate RMSE
rmse = np.sqrt(mean_squared_error(test_df['Prices'], predictions))

# Calculate MAE
mae = mean_absolute_error(test_df['Prices'], predictions)

# Print the evaluation metrics
print(f'RMSE: {rmse:.3f}')
print(f'MAE: {mae:.3f}')

RMSE: 0.288
MAE: 0.208

```

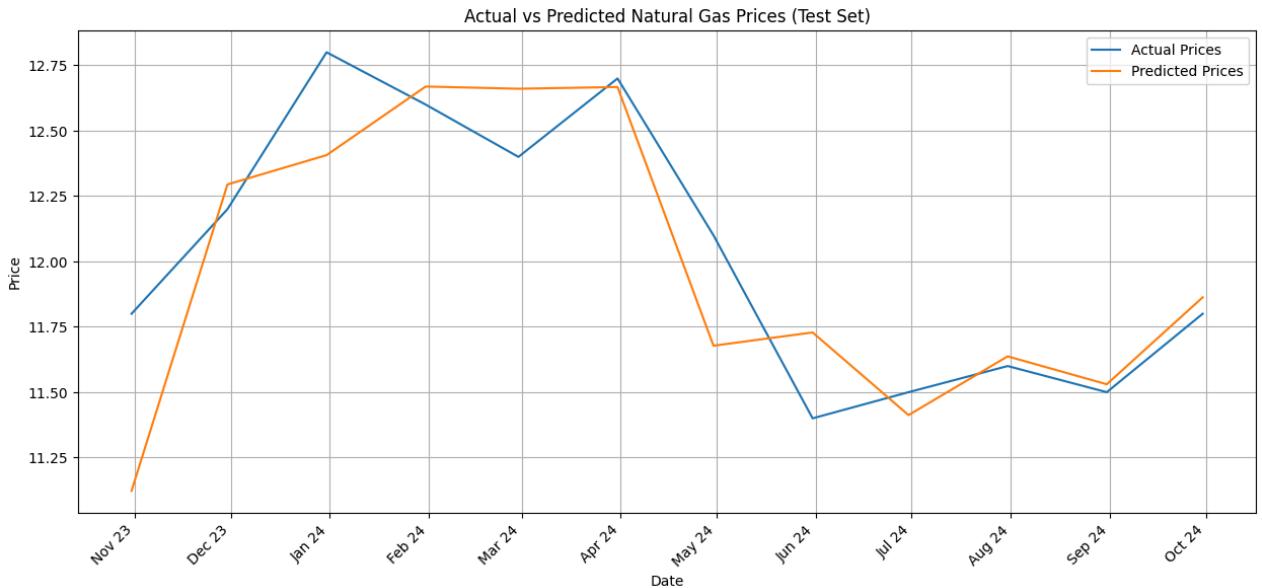
Start coding or [generate](#) with AI.

```

plt.figure(figsize=(15, 7))
plt.plot(test_df['Dates'], test_df['Prices'], label='Actual Prices')
plt.plot(test_df['Dates'], predictions, label='Predicted Prices')

plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Actual vs Predicted Natural Gas Prices (Test Set)')
plt.legend()
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=1))
plt.gcf().autofmt_xdate(rotation=45)
plt.grid(True)
plt.show()

```



```

plt.figure(figsize=(15, 7))

# Plot historical data by year and month with year labels
for year in df['Year'].unique():
    yearly_data = df[df['Year'] == year]
    plt.plot(yearly_data['Month'], yearly_data['Prices'], label=year)
    # Add year label at the starting point of each line
    first_month_data = yearly_data[yearly_data['Month'] == yearly_data['Month'].min()]
    if not first_month_data.empty:
        plt.text(first_month_data['Month'].iloc[0], first_month_data['Prices'].iloc[0], str(year), fontsize=9, ha='right')

# Plot the forecasted values
# We need to align the forecast with the months after the last historical data point
last_month_historical = df['Month'].iloc[-1]
start_month_forecast = last_month_historical + 1
if start_month_forecast > 12:
    start_month_forecast = start_month_forecast - 12

forecast_months = [(last_month_historical + i - 1) % 12 + 1 for i in range(1, forecast_periods + 1)]

# Plot only the forecasted mean, not the confidence intervals
plt.plot(forecast_months, forecast, label='Forecasted Prices', color='red', linestyle='--')

# Add year labels for the forecasted line (assuming 2024 and 2025 based on the data)

```

```

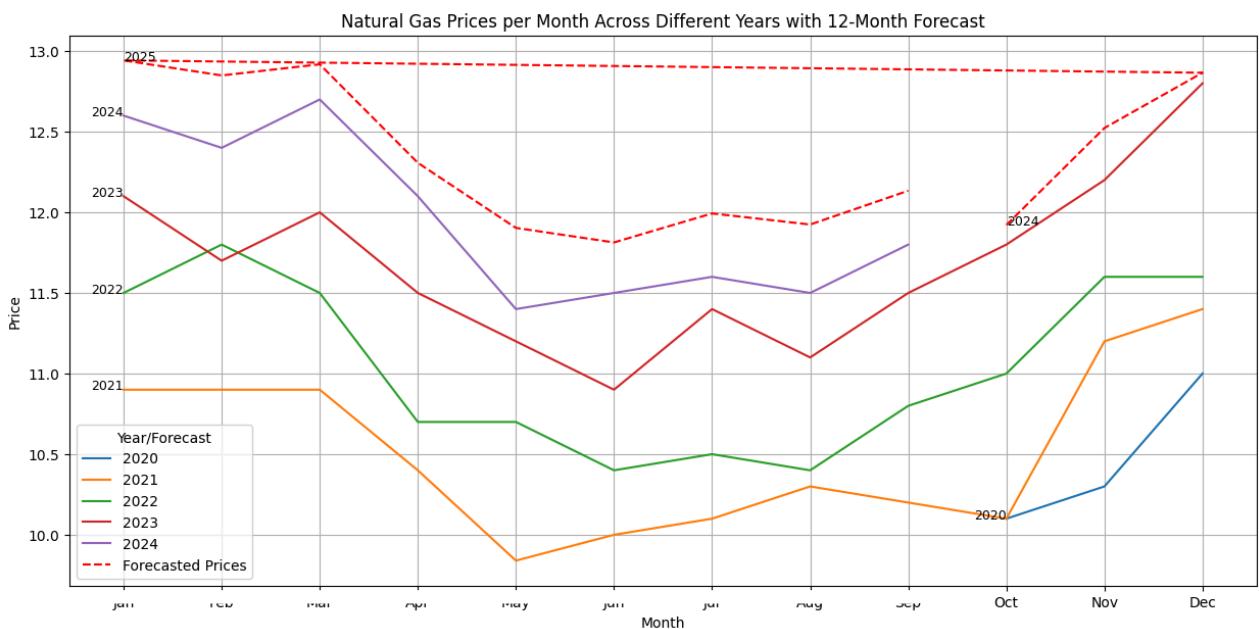
# We need to determine the correct year for each forecast point based on the last historical date
last_date = df['Dates'].iloc[-1]
forecast_dates = pd.date_range(start=last_date, periods=forecast_periods + 1, freq='M')[1:]

# Find the index where the year changes from 2024 to 2025 in the forecast dates
year_change_index = None
for i in range(len(forecast_dates) - 1):
    if forecast_dates[i].year != forecast_dates[i+1].year:
        year_change_index = i + 1
        break

if not forecast_dates.empty:
    plt.text(forecast_dates[0].month, forecast.iloc[0], str(forecast_dates[0].year), fontsize=9, ha='left')
    if year_change_index is not None:
        plt.text(forecast_dates[year_change_index].month, forecast.iloc[year_change_index], str(forecast_dates[year_change_index].year), fontsize=9, ha='left')

plt.xlabel('Month')
plt.ylabel('Price')
plt.title('Natural Gas Prices per Month Across Different Years with 12-Month Forecast')
plt.xticks(range(1, 13), ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'])
plt.legend(title='Year/Forecast')
plt.grid(True)
plt.show()

```



Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```

from statsmodels.tsa.holtwinters import ExponentialSmoothing

# Instantiate the Holt-Winters model
# Using additive trend and additive seasonality with a seasonal period of 12
model_holt = ExponentialSmoothing(df['Prices'], seasonal_periods=12, trend='add', seasonal='add')

# Fit the model
model_holt_fit = model_holt.fit()

# Summarize model_fit
display(model_holt_fit.summary())

```

```

ExponentialSmoothing Model Results
Dep. Variable: Prices No. Observations: 48
Model: ExponentialSmoothing SSE 1.212
Optimized: True AIC -144.574
Trend: Additive BIC -114.635
Seasonal: Additive AICC -120.988
Seasonal Periods: 12 Date: Tue, 02 Dec 2025
Box-Cox: False Time: 06:54:13
Box-Cox Coeff.: None

      coeff  code optimized
smoothing_level 1.4901e-08 alpha  True
smoothing_trend 1.4737e-10 beta  True
smoothing_seasonal 0.000000 gamma True
initial_level    10.121153 l.0   True
initial_trend     0.0451251 b.0   True
initial_seasons.0 -0.2285319 s.0  True
initial_seasons.1 0.3013445 s.1  True
initial_seasons.2 0.6312183 s.2  True
initial_seasons.3 0.6610935 s.3  True
initial_seasons.4 0.5409683 s.4  True
initial_seasons.5 0.5708434 s.5  True
initial_seasons.6 -0.0742821 s.6  True
initial_seasons.7 -0.5094058 s.7  True
initial_seasons.8 -0.6395313 s.8  True
initial_seasons.9 -0.4846558 s.9  True
initial_seasons.10 -0.6047806 s.10 True
initial_seasons.11 -0.3999058 s.11 True

```

```

# Generate the 12-month forecast
holt_forecast = model_holt_fit.forecast(steps=12)

# Display the forecast
print("Holt-Winters 12-Month Forecast:")
display(holt_forecast)

```

Holt-Winters 12-Month Forecast:

	0
48	12.103751
49	12.678752
50	13.053751
51	13.128751
52	13.053751
53	13.128751
54	12.528751
55	12.138752
56	12.053752
57	12.253753
58	12.178753
59	12.428753

dtype: float64

Start coding or [generate](#) with AI.

```

# Generate predictions for the testing period using the fitted Holt-Winters model
holt_predictions = model_holt_fit.predict(start=test_df.index[0], end=test_df.index[-1])

# Calculate RMSE for Holt-Winters
rmse_holt = np.sqrt(mean_squared_error(test_df['Prices'], holt_predictions))

# Calculate MAE for Holt-Winters
mae_holt = mean_absolute_error(test_df['Prices'], holt_predictions)

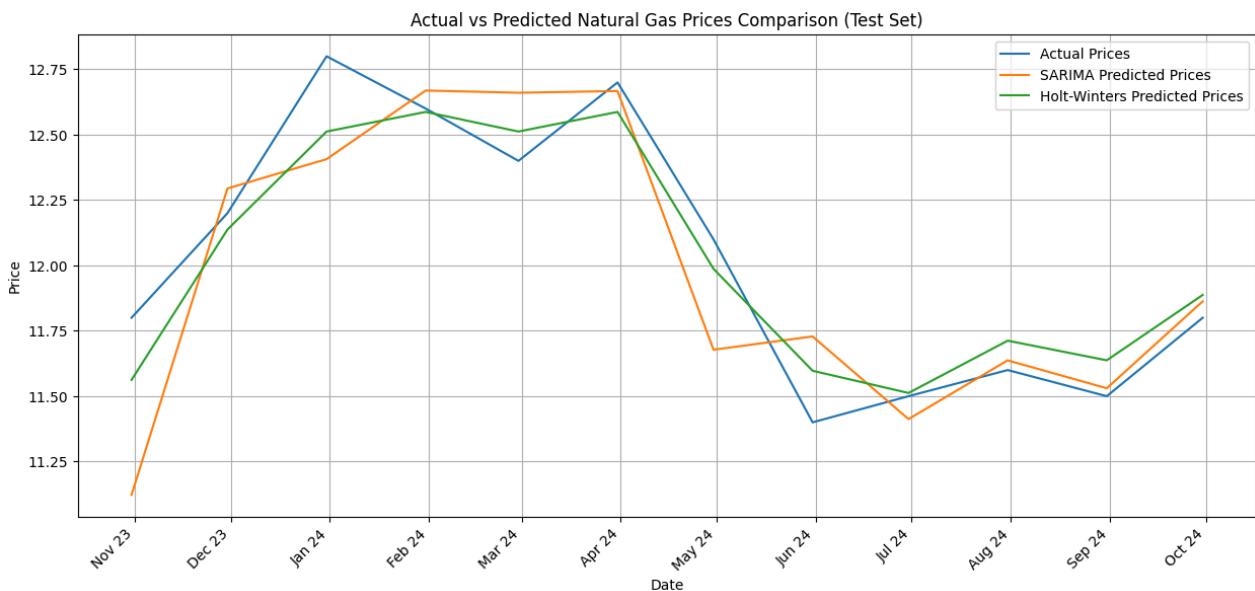
# Print the evaluation metrics for Holt-Winters
print(f'Holt-Winters RMSE on Test Set: {rmse_holt:.3f}')
print(f'Holt-Winters MAE on Test Set: {mae_holt:.3f}')

```

```
Holt-Winters RMSE on Test Set: 0.147
Holt-Winters MAE on Test Set: 0.124
```

```
plt.figure(figsize=(15, 7))
plt.plot(test_df['Dates'], test_df['Prices'], label='Actual Prices')
plt.plot(test_df['Dates'], predictions, label='SARIMA Predicted Prices')
# plt.plot(test_df['Dates'], prophet_predictions, label='Prophet Predicted Prices') # Removed Prophet line
plt.plot(test_df['Dates'], holt_predictions, label='Holt-Winters Predicted Prices')

plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Actual vs Predicted Natural Gas Prices Comparison (Test Set)')
plt.legend()
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=1))
plt.gcf().autofmt_xdate(rotation=45)
plt.grid(True)
plt.show()
```



Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```
plt.figure(figsize=(15, 7))

# Plot historical data
plt.plot(df['Dates'], df['Prices'], label='Historical Prices')

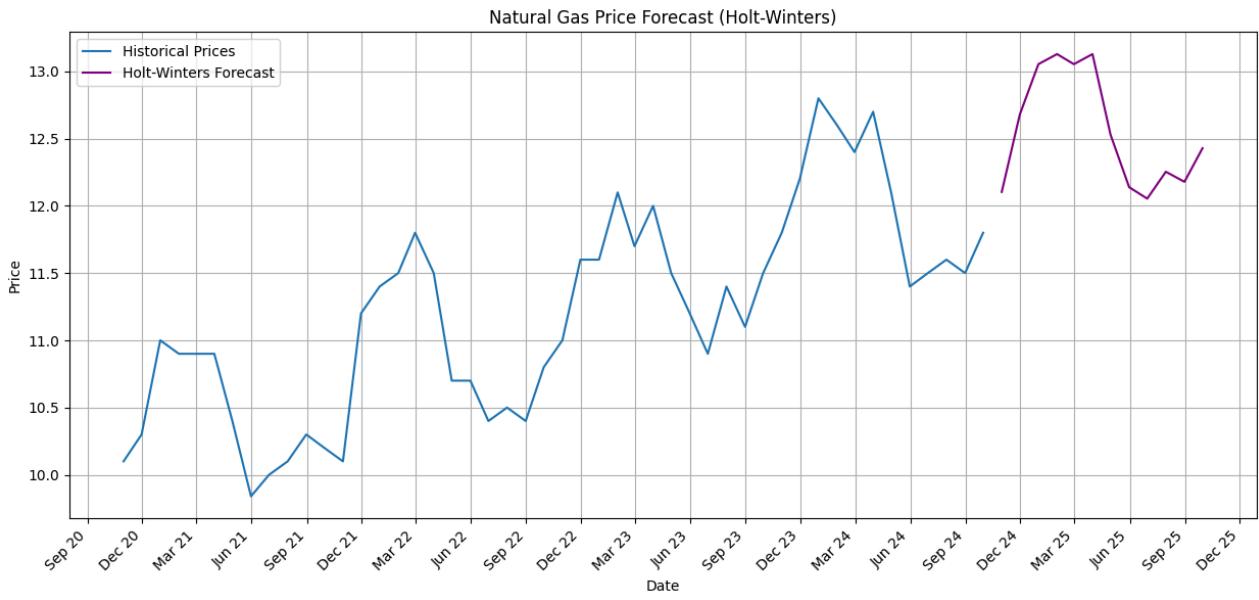
# Create a date range for the Holt-Winters forecast
last_date = df['Dates'].iloc[-1]
holt_forecast_dates = pd.date_range(start=last_date, periods=12 + 1, freq='M')[1:]

# Plot the Holt-Winters forecasted values
plt.plot(holt_forecast_dates, holt_forecast, label='Holt-Winters Forecast', color='purple')

plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Natural Gas Price Forecast (Holt-Winters)')
plt.legend()

# Improve date formatting
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=3)) # Show every 3 months for clarity
plt.gcf().autofmt_xdate(rotation=45)

plt.grid(True)
plt.show()
```



Start coding or [generate](#) with AI.

```
plt.figure(figsize=(15, 7))

# Plot historical data
plt.plot(df['Dates'], df['Prices'], label='Historical Prices')

# Create date ranges for the forecasts
last_date = df['Dates'].iloc[-1]
forecast_dates = pd.date_range(start=last_date, periods=12 + 1, freq='M')[1:]

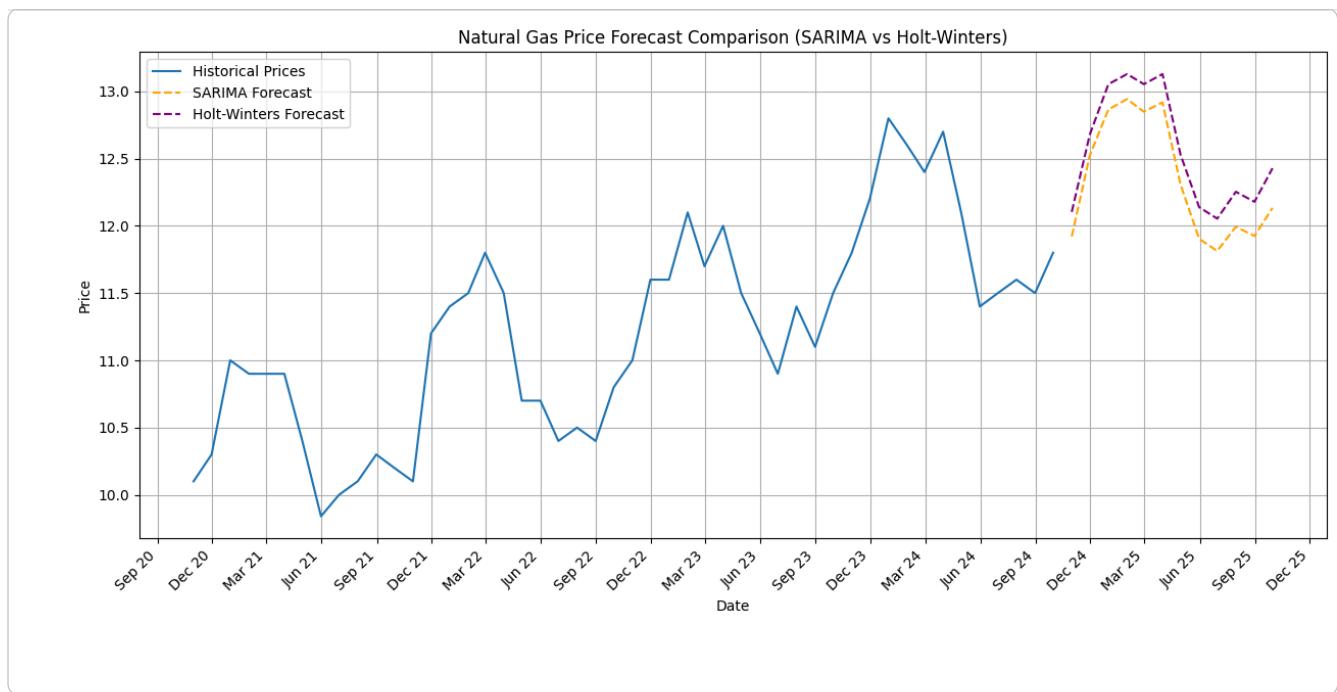
# Plot SARIMA forecast
plt.plot(forecast_dates, forecast, label='SARIMA Forecast', color='orange', linestyle='--')

# Plot Holt-Winters forecast
plt.plot(forecast_dates, holt_forecast, label='Holt-Winters Forecast', color='purple', linestyle='--')

plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Natural Gas Price Forecast Comparison (SARIMA vs Holt-Winters)')
plt.legend()

# Improve date formatting
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=3)) # Show every 3 months for clarity
plt.gcf().autofmt_xdate(rotation=45)

plt.grid(True)
plt.show()
```



Start coding or [generate](#) with AI.

```
#!pip install prophet
```

Prophet requires the dataframe to have two columns: `ds` for the time series (datetime format) and `y` for the measurements.

```
# Prepare data for Prophet
prophet_df = df[['Dates', 'Prices']].rename(columns={'Dates': 'ds', 'Prices': 'y'})
display(prophet_df.head())
```

	ds	y
0	2020-10-31	10.1
1	2020-11-30	10.3
2	2020-12-31	11.0
3	2021-01-31	10.9
4	2021-02-28	10.9

Train the Prophet model on the entire dataset. Prophet automatically handles seasonality.

```
from prophet import Prophet

# Instantiate the Prophet model
model_prophet = Prophet()

# Fit the model to the entire dataset
model_prophet.fit(prophet_df)

INFO:prophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to override this.
INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.
<prophet.forecaster.Prophet at 0x7a9abb4b5520>
```

Generate the 12-month forecast using the trained Prophet model.

```
# Create a dataframe with future dates for the next 12 months
future_prophet = model_prophet.make_future_dataframe(periods=12, freq='M')

# Generate predictions for the future dates
forecast_prophet_future = model_prophet.predict(future_prophet)

# Extract the future forecast (excluding the historical data)
prophet_future_predictions = forecast_prophet_future[forecast_prophet_future['ds'] > df['Dates'].max()]

# Display the forecasted values in a table
print("Prophet Forecasted prices for the next 12 months:")
display(prophet_future_predictions[['ds', 'yhat']])
```

Prophet Forecasted prices for the next 12 months:

	ds	yhat
48	2024-10-31	11.960243
49	2024-11-30	12.497213
50	2024-12-31	12.915378
51	2025-01-31	12.939208
52	2025-02-28	12.875277
53	2025-03-31	12.951606
54	2025-04-30	12.339298
55	2025-05-31	11.853790
56	2025-06-30	11.848568
57	2025-07-31	11.983318
58	2025-08-31	11.949649
59	2025-09-30	12.153201

Visualize the historical data and the 12-month Prophet forecast.

```
plt.figure(figsize=(15, 7))

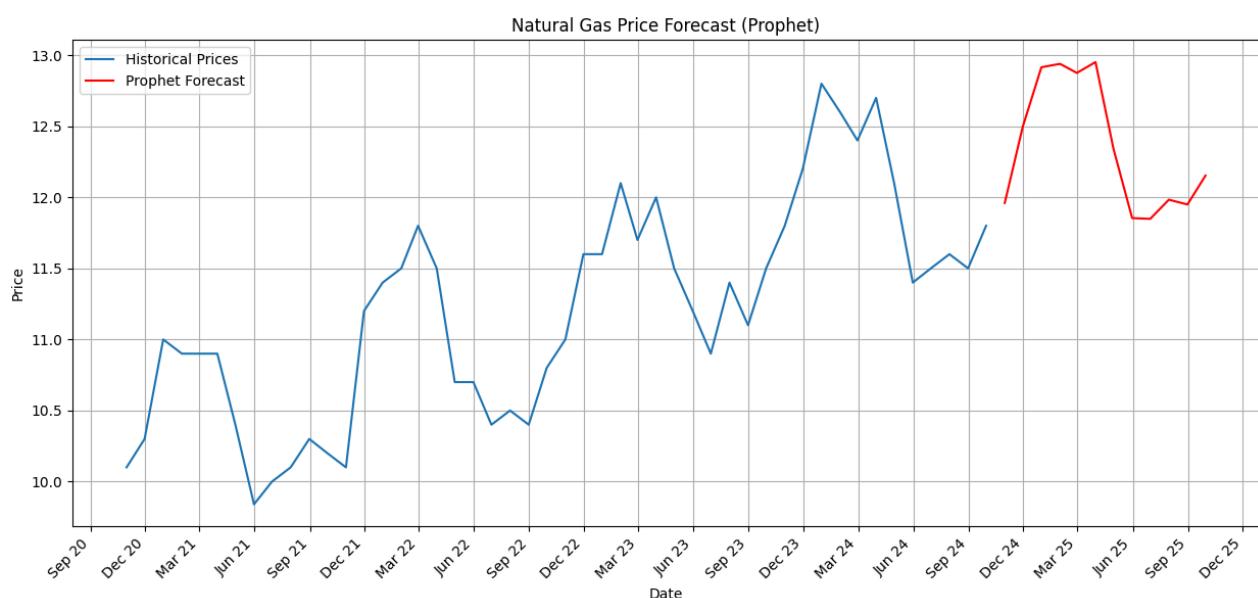
# Plot historical data
plt.plot(df['Dates'], df['Prices'], label='Historical Prices')

# Plot the Prophet forecasted values for the next 12 months
plt.plot(prophet_future_predictions['ds'], prophet_future_predictions['yhat'], label='Prophet Forecast', color='red')

plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Natural Gas Price Forecast (Prophet)')
plt.legend()

# Improve date formatting
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=3)) # Show every 3 months for clarity
plt.gcf().autofmt_xdate(rotation=45)

plt.grid(True)
plt.show()
```



```
from sklearn.metrics import mean_squared_error, mean_absolute_error
from prophet import Prophet # Ensure Prophet is imported

# Split prophet_df into training and testing sets
# test_size is already defined in the notebook context (e.g., test_size = 12)
prophet_train_df = prophet_df[:-test_size]
```

```

prophet_test_df = prophet_df[-test_size:]

# Instantiate and fit a new Prophet model on the training data for evaluation
# This ensures that the model has not seen the test data during training for performance metrics
model_prophet_eval = Prophet()
model_prophet_eval.fit(prophet_train_df)

# Create a future dataframe for the test period (ds column only)
future_prophet_test = prophet_test_df[['ds']]

# Generate predictions for the test period using the evaluation model
forecast_prophet_test = model_prophet_eval.predict(future_prophet_test)
prophet_predictions = forecast_prophet_test['yhat']

# Calculate RMSE for Prophet
rmse_prophet = np.sqrt(mean_squared_error(prophet_test_df['y'], prophet_predictions))

# Calculate MAE for Prophet
mae_prophet = mean_absolute_error(prophet_test_df['y'], prophet_predictions)

# Print the evaluation metrics for Prophet
print(f'Prophet RMSE on Test Set: {rmse_prophet:.3f}')
print(f'Prophet MAE on Test Set: {mae_prophet:.3f}')

INFO:prophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to override this.
INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.
Prophet RMSE on Test Set: 0.399
Prophet MAE on Test Set: 0.336

```

```

plt.figure(figsize=(15, 7))

# Plot historical data
plt.plot(df['Dates'], df['Prices'], label='Historical Prices')

# Create date ranges for the forecasts
last_date = df['Dates'].iloc[-1]
forecast_dates = pd.date_range(start=last_date, periods=12 + 1, freq='M')[1:]

# Plot SARIMA forecast
plt.plot(forecast_dates, forecast, label='SARIMA Forecast', color='orange', linestyle='--')

# Plot Holt-Winters forecast
plt.plot(forecast_dates, holt_forecast, label='Holt-Winters Forecast', color='purple', linestyle='--')

# Plot Prophet forecast
# Ensure prophet_future_predictions is available from previous steps or regenerate it
# Assuming prophet_future_predictions is available and contains the forecast for the next 12 months
plt.plot(prophet_future_predictions['ds'], prophet_future_predictions['yhat'], label='Prophet Forecast', color='green')

plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Natural Gas Price Forecast Comparison (Historical + Models)')
plt.legend()

# Improve date formatting
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=3)) # Show every 3 months for clarity
plt.gcf().autofmt_xdate(rotation=45)

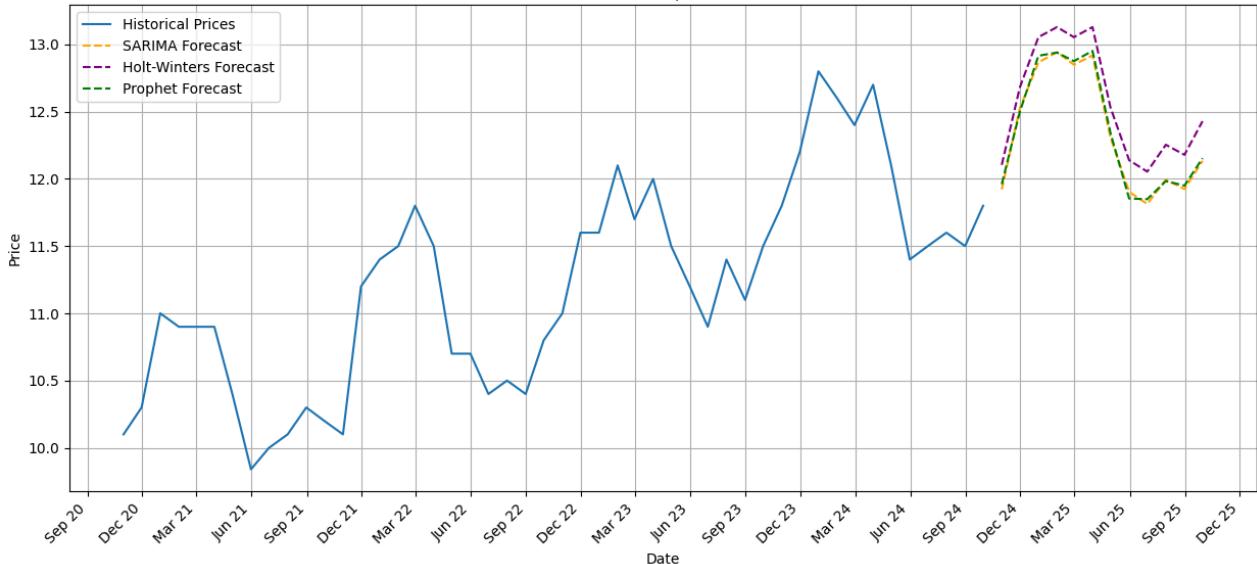
plt.grid(True)

# Save the figure as a PNG file
plt.savefig('natural_gas_price_forecast_comparison.png')

plt.show()

```

Natural Gas Price Forecast Comparison (Historical + Models)



```
import pandas as pd

# Create a dictionary to store the metrics
metrics_data = {
    'Model': ['SARIMA', 'Prophet', 'Holt-Winters'],
    'RMSE': [rmse, rmse_prophet, rmse_holt],
    'MAE': [mae, mae_prophet, mae_holt]
}

# Create a pandas DataFrame from the dictionary
metrics_df = pd.DataFrame(metrics_data)

# Display the table
print("Model Performance Comparison on Test Set:")
display(metrics_df)
```

Model Performance Comparison on Test Set:

	Model	RMSE	MAE	Actions
0	SARIMA	0.288160	0.208036	
1	Prophet	0.399289	0.335605	
2	Holt-Winters	0.147071	0.123751	

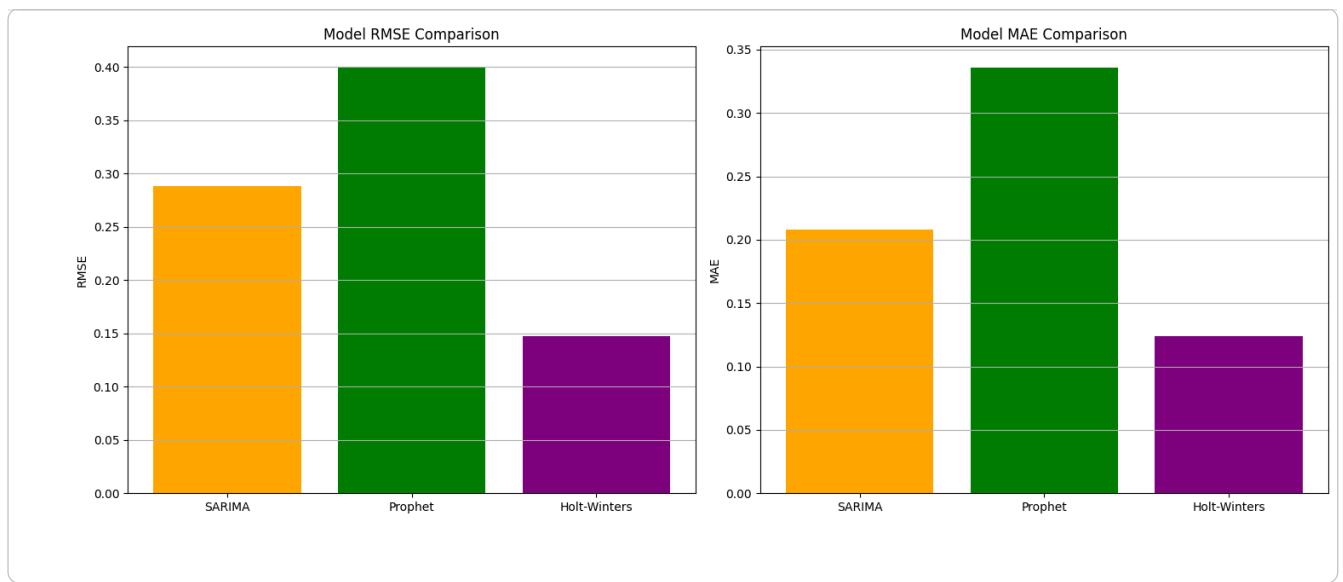
Next steps: [Generate code with metrics_df](#) [New interactive sheet](#)

```
# Visualize the comparison of evaluation metrics
fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Plot RMSE
axes[0].bar(metrics_df['Model'], metrics_df['RMSE'], color=['orange', 'green', 'purple'])
axes[0].set_ylabel('RMSE')
axes[0].set_title('Model RMSE Comparison')
axes[0].grid(axis='y')

# Plot MAE
axes[1].bar(metrics_df['Model'], metrics_df['MAE'], color=['orange', 'green', 'purple'])
axes[1].set_ylabel('MAE')
axes[1].set_title('Model MAE Comparison')
axes[1].grid(axis='y')

plt.tight_layout()
plt.show()
```



```

import ipywidgets as widgets
from IPython.display import display
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.dates as mdates

# Get the last historical date and price
last_historical_date = df['Dates'].iloc[-1]
last_historical_price = df['Prices'].iloc[-1]

# Define the specific historical date range to display (Oct 2022 to Oct 2024)
historical_start_date = pd.to_datetime('2022-10-31')
historical_end_date = pd.to_datetime('2024-10-31')

# Filter the DataFrame for the specified historical range
historical_display_df = df[(df['Dates'] >= historical_start_date) & (df['Dates'] <= historical_end_date)].copy()

# Define the specific future dates for selection
future_dates_options = {
    '11 Aug 2026': pd.to_datetime('2026-08-11'),
    '26 Jan 2027': pd.to_datetime('2027-01-26'),
    '30 Oct 2050': pd.to_datetime('2050-10-30') # Changed one option to 30 Oct 2050
}

# Create a dropdown widget for predefined dates
dropdown_dates = widgets.Dropdown(
    options=[(key, value) for key, value in future_dates_options.items()],
    description='Choose a Future Date:',
    layout=widgets.Layout(width='50%') # Adjust layout for better visibility
)

# Create a date picker widget for custom date input
date_picker = widgets.DatePicker(
    description='Or Pick a Specific Date:',
    disabled=False,
    layout=widgets.Layout(width='50%') # Adjust layout for better visibility
)

# Create a dropdown for model selection
model_options = ['Holt-Winters (Best)', 'SARIMA', 'Prophet']
dropdown_models = widgets.Dropdown(
    options=model_options,
    description='Choose Model:',
    layout=widgets.Layout(width='50%') # Adjust layout for better visibility
)

# Create an output widget to display the prediction
output = widgets.Output()

# Function to make prediction and visualize
def predict_and_plot(date, model_name):
    with output:
        output.clear_output()
        if date:
            # Ensure the selected date is a datetime object
            if isinstance(date, str):
                date = pd.to_datetime(date)

```

```

# Calculate the number of steps ahead from the last historical date
steps_ahead = (date.year - last_historical_date.year) * 12 + (date.month - last_historical_date.month)
# Adjust steps_ahead if the day of the month is before the last historical date's day in the same month
if date.year == last_historical_date.year and date.month == last_historical_date.month and date.day < last_historical_date.day:
    steps_ahead = 0 # Do not predict for dates on or before the last historical date
elif (date.year == last_historical_date.year and date.month < last_historical_date.month) or date.year < last_historical_date.year:
    steps_ahead = 0 # Do not predict for past dates

if steps_ahead > 0:
    # Select the model based on user choice
    if model_name == 'SARIMA':
        # Ensure SARIMA model is fitted on the entire data if not already
        try:
            model_final_fit # Check if it exists
        except NameError:
            print("SARIMA model not fitted on entire data. Please run the SARIMA training cell.")
            return
        predicted_price = model_final_fit.forecast(steps=steps_ahead).iloc[-1]
        used_model_name = 'SARIMA'
    elif model_name == 'Prophet':
        # Ensure Prophet model is fitted on the entire data if not already
        try:
            model_prophet # Check if it exists
        except NameError:
            print("Prophet model not fitted. Please run the Prophet training cell.")
            return
        # Prophet requires a future dataframe with the specific date
        future_df_single = pd.DataFrame({'ds': [date]})
        predicted_price = model_prophet.predict(future_df_single)['yhat'].iloc[0]
        used_model_name = 'Prophet'
    else: # Default to Holt-Winters or if Holt-Winters is explicitly chosen
        # Ensure Holt-Winters model is fitted on the entire data if not already
        try:
            model_holt_fit # Check if it exists
        except NameError:
            print("Holt-Winters model not fitted. Please run the Holt-Winters training cell.")
            return
        predicted_price = model_holt_fit.forecast(steps=steps_ahead).iloc[-1]
        used_model_name = 'Holt-Winters'

print(f"Predicted price for {date.strftime('%Y-%m-%d')}: {predicted_price:.3f} (using {used_model_name})")

# Prepare data for the comparison bar chart
# Combine historical data within the specified range and the predicted price
plot_dates = historical_display_df['Dates'].tolist() + [date]
plot_prices = historical_display_df['Prices'].tolist() + [predicted_price]
plot_colors = ['gray'] * len(historical_display_df) + ['skyblue'] # Colors for historical and predicted

plt.figure(figsize=(15, 6))
plt.bar(plot_dates, plot_prices, color=plot_colors, width=20) # Use dates as x, adjust width

plt.ylabel('Price')
plt.title(f'Natural Gas Price: Historical (Oct 2022 - Oct 2024) vs Predicted ({used_model_name} Model)')

# Format x-axis as dates, skipping dates between Oct 2024 and input date
ax = plt.gca()
ax.xaxis.set_major_formatter(mdates.DateFormatter('%b %y'))
ax.xaxis.set_major_locator(mdates.MonthLocator(interval=3)) # Show every 3 months for historical

# Adjust x-axis limits to show a gap and include the predicted date
# This is a bit tricky with standard matplotlib bar charts on date axis
# A workaround is to set explicit tick locations and labels
tick_dates = historical_display_df['Dates'].tolist() + [date]
ax.set_xticks(tick_dates)
ax.set_xticklabels([d.strftime('%b %y') for d in tick_dates], rotation=45, ha='right')

plt.tight_layout()
plt.show()

else:
    print("Please select a date after the last historical date.")

# Function to handle date selection change
def handle_date_change(change):
    selected_date = change['new']
    selected_model = dropdown_models.value # Get the currently selected model
    predict_and_plot(selected_date, selected_model)

```

```

predict_and_plot(selected_date, selected_model)

# Function to handle model selection change
def handle_model_change(change):
    selected_model = change['new']
    # Trigger prediction with the currently selected date
    # Check which date picker has a value
    if date_picker.value:
        selected_date = date_picker.value
    elif dropdown_dates.value:
        selected_date = dropdown_dates.value
    else:
        selected_date = None # No date selected yet

    if selected_date:
        predict_and_plot(selected_date, selected_model)

# Link the dropdowns and date picker to the prediction functions
dropdown_dates.observe(handle_date_change, names='value')
date_picker.observe(handle_date_change, names='value')
dropdown_models.observe(handle_model_change, names='value')

# Display the widgets
print("Pick a Date and Model and I will show you the predicted price:")
display(dropdown_dates, date_picker, dropdown_models, output)

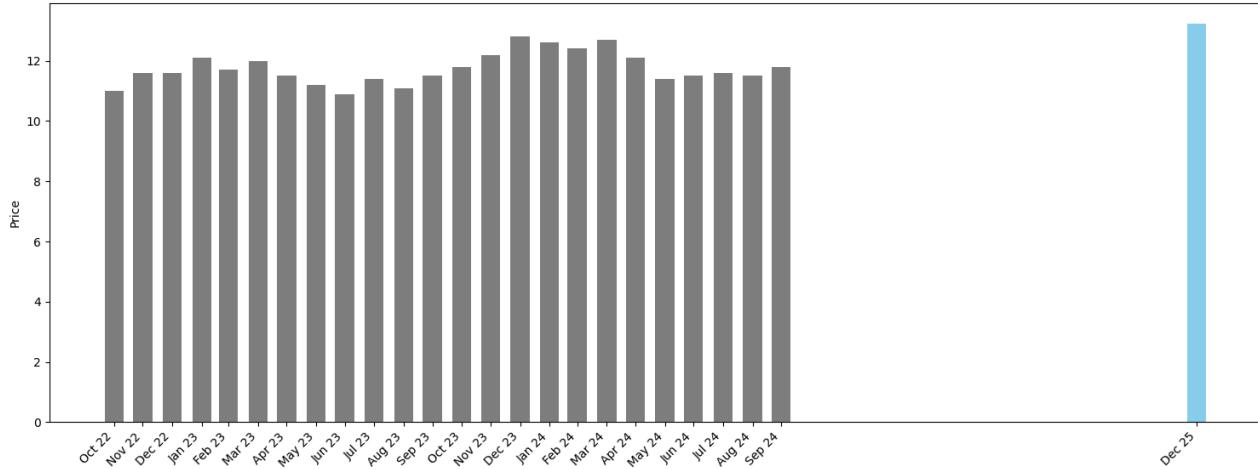
```

Pick a Date and Model and I will show you the predicted price:

Choose a F...	26 Jan 2027
Or Pick a ...	10-12-2025
Choose Mo...	SARIMA

Predicted price for 2025-12-10: 13.239 (using SARIMA model)

Natural Gas Price: Historical (Oct 2022 - Oct 2024) vs Predicted (SARIMA Model)



Natural Gas Price Forecasting Project Report

Project Overview

This project aimed to analyze historical natural gas prices, identify patterns including seasonality, and develop forecasting models to predict future prices. We explored the data, handled necessary transformations, implemented and evaluated several forecasting models, and created an interactive tool for future price predictions.

Data Analysis and Key Findings

- Data Loading and Exploration:** The natural gas price data was successfully loaded and explored, providing an initial understanding of its structure and content. The dataset contains 48 entries with two columns: 'Dates' (initially object type) and 'Prices' (float64).
- Data Preprocessing:** The 'Dates' column was converted to datetime objects, which is essential for time series analysis.